

# KUBERNETES

## What is kubernetes

Kubernetes (popularly known as k8s) is a container orchestration engine, which as the name suggests lets you create, delete, and update containers

This is useful when:

You have your docker images in the docker registry and want to deploy it in a cloud native fashion

You want to not worry about patching, crashes. You want the system to auto heal

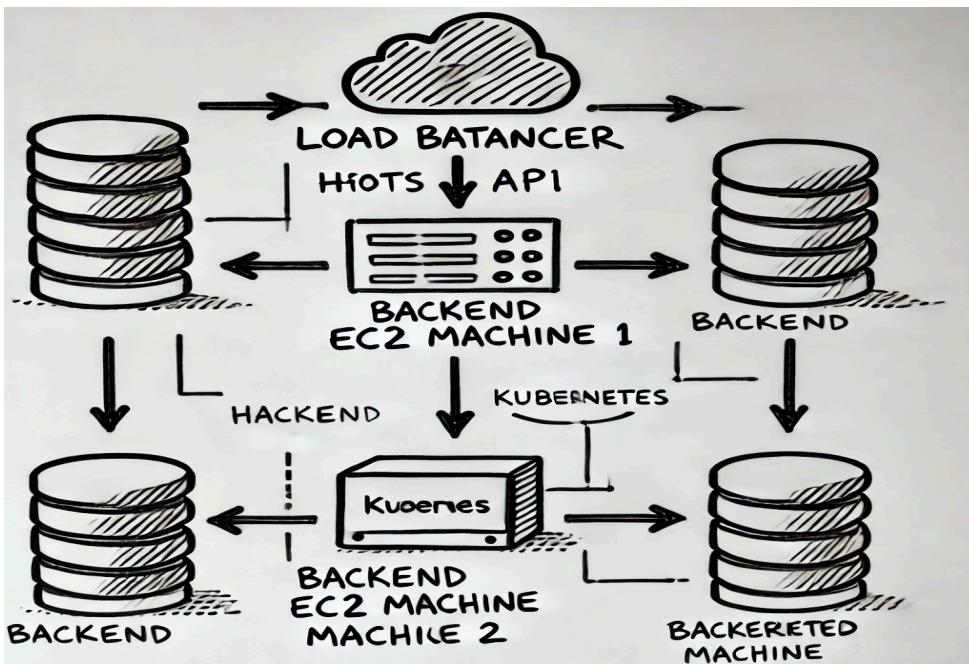
You want to autoscale with some simple constructs

You want to observe your complete system in a simple dashboard

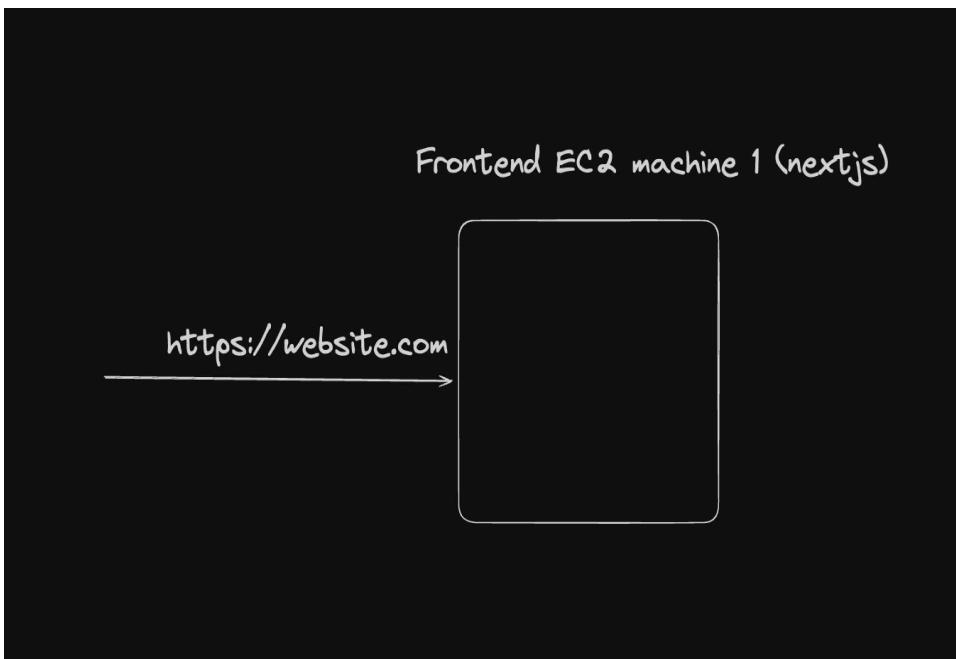


# Before kubernetes

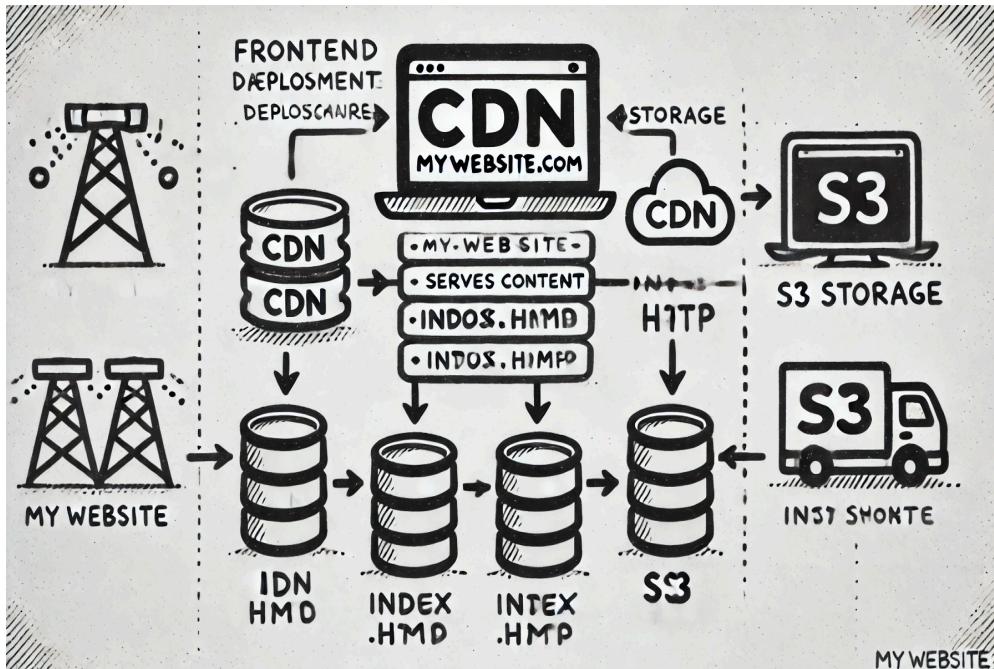
## Backend



## Frontend (Nextjs)

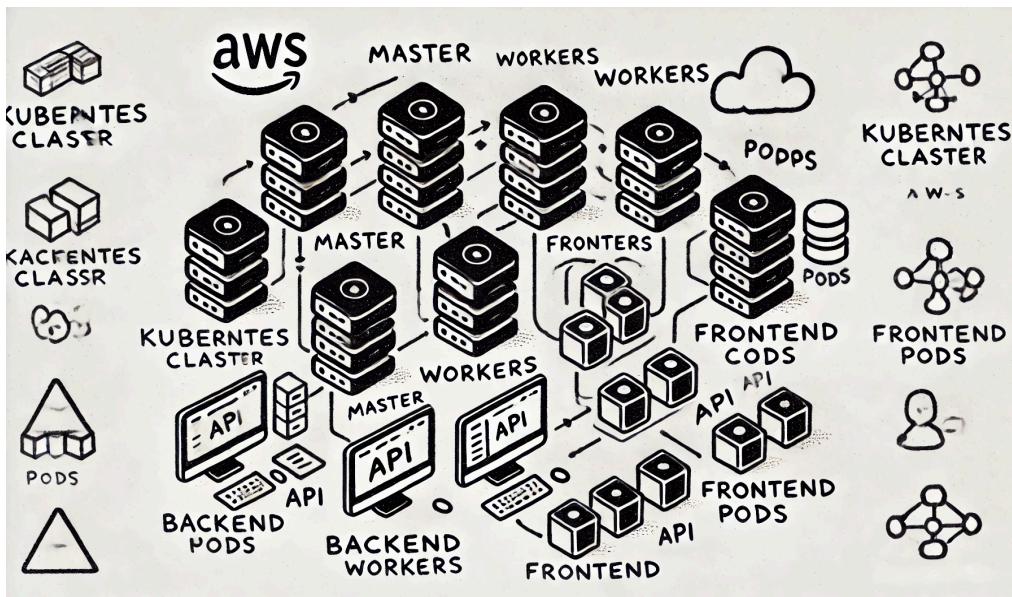


# Frontend (React)



## After kubernetes

Your frontend, backend are all **pods** in your **kubernetes cluster**



## Nodes:-

In kubernetes, you can create and connect various machines together, all of which are running kubernetes. Every machine here is known as a node

There are two types of nodes:

a. Master Node (Control pane) - The node that takes care of deploying the containers/healing them/listening to the developer to understand what to deploy

b. Worker Nodes - The nodes that actually run your Backend/frontend

## Cluster

A bunch of worker nodes + master nodes make up your kubernetes cluster. You can always add more / remove nodes from a cluster.

## Images:-

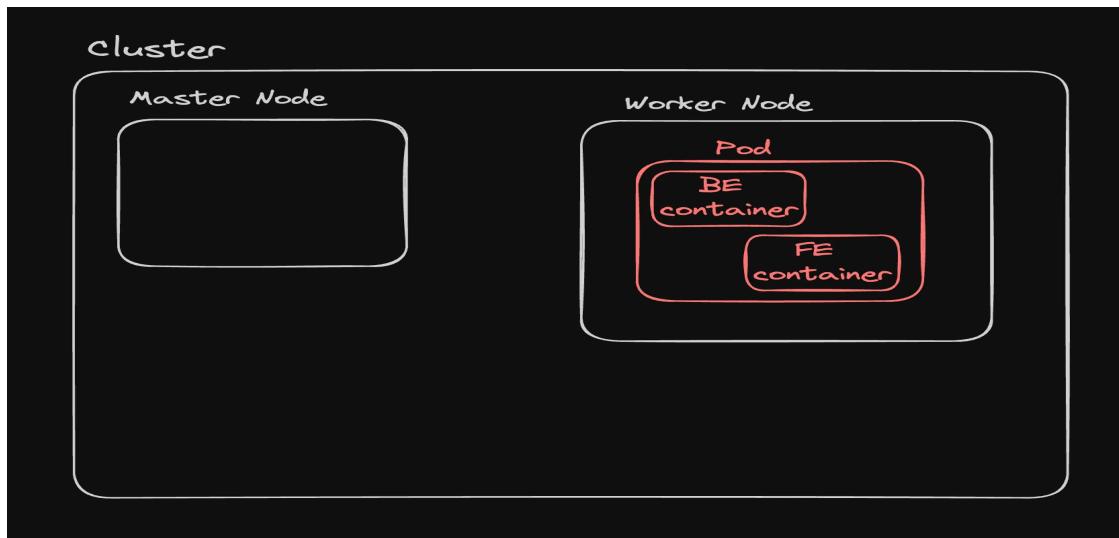
A Docker image is a lightweight, standalone, and executable software package that includes everything needed to run a piece of software, including the code, runtime, libraries, environment variables, and configuration files. Images are built from a set of instructions defined in a file called a Dockerfile.

## Containers

A container is an image in execution. For example if you run

### Pods:-

A pod is the smallest and simplest unit in the Kubernetes object model that you can create or deploy



## **Creating a k8s cluster**

Locally (Make sure you have docker):-

minikube

kind

On cloud:-

GKE, AWS K8s, vultr

Single node setup:-

Create a 1 node cluster

```
kind create cluster --name local
```

Check the docker containers you have running

```
docker ps
```

You will notice a single container running (control-plane)

Delete the cluster

```
kind delete cluster -n local
```

## Multi node setup

Create a `clusters.yml` file

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
- role: worker
- role: worker
```

Create the node setup

```
kind create cluster --config clusters.yml --name local
```

Check docker containers

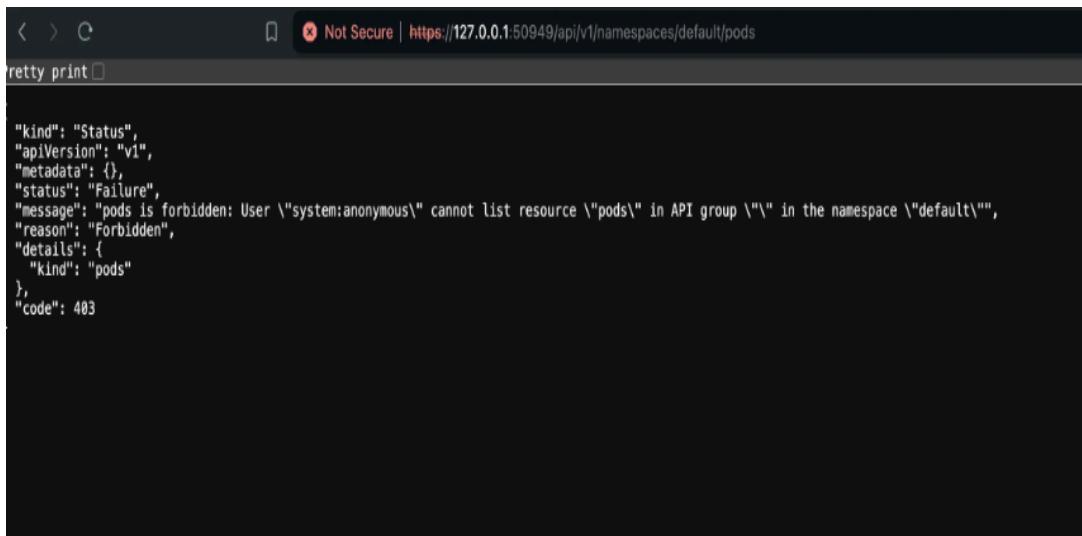
```
docker ps
```

## Kubernetes API

The master node (control pane) exposes an API that a developer can use to start pods.

Try the API

Run docker ps to find where the control pane is running



A screenshot of a web browser window. The address bar shows a URL starting with "https://127.0.0.1:50949/api/v1/namespaces/default/pods". Below the address bar, there is a "pretty print" button. The main content area displays a JSON error response:

```
{"kind": "Status", "apiVersion": "v1", "metadata": {}, "status": "Failure", "message": "pods is forbidden: User \"system:anonymous\" cannot list resource \"pods\" in API group \"\" in the namespace \"default\"", "reason": "Forbidden", "details": { "kind": "pods" }, "code": 403}
```

Kubernetes API server does authentication checks and prevents you from getting in.

All of your authorization credentials are stored by kind in

`~/.kube/config`

# kubectl

`kubectl` is a command-line tool for interacting with Kubernetes clusters. It provides a way to communicate with the Kubernetes API server and manage Kubernetes resources.

## Ensure kubectl works fine

```
kubectl get nodes  
kubectl get pods
```

If you want to see the exact HTTP request that goes out to the API server, you can add `--v=8` flag

```
kubectl get nodes --v=8
```

# Creating a Pod

There were 5 jargons we learnt about:-

Cluster

Nodes

Images

Containers

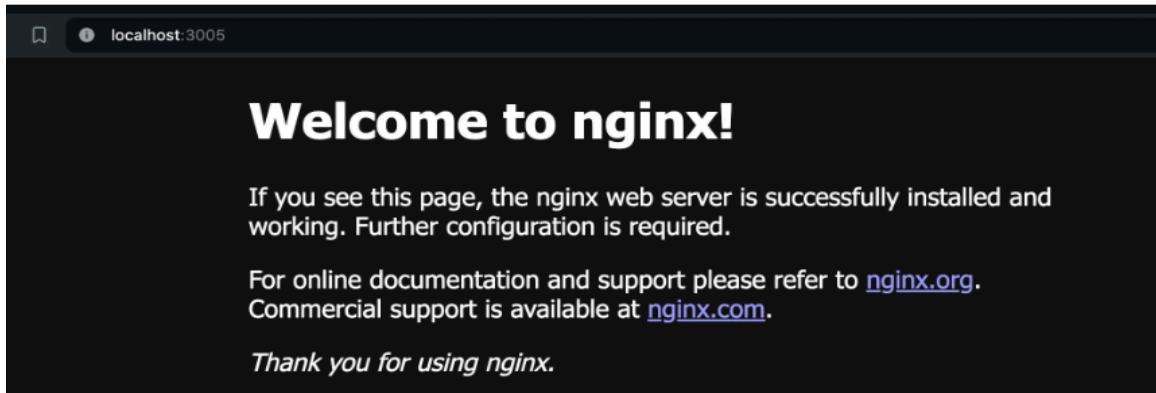
Pods

## Starting using docker

```
docker run -p 3005:80 nginx
```

```
docker (com.docker.cli)
~ $ docker run -p 3005:80 nginx
Unable to find image 'nginx:latest' locally
latest: Pulling from library/nginx
24c63b8dcbb6: Already exists
cfe05c174c17: Pull complete
f72850762d9b: Pull complete
82945d645e12: Pull complete
a6f455b832da: Pull complete
22fc067a16f6: Pull complete
4bd9fc5c1e9f: Pull complete
Digest: sha256:0f04e4f646a3f14bf31d8bc8d885b6c951fdcf42589d06845f64d18aec6a3c4d
Status: Downloaded newer image for nginx:latest
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/06/01 00:33:21 [notice] 1#1: using the "epoll" event method
2024/06/01 00:33:21 [notice] 1#1: nginx/1.27.0
2024/06/01 00:33:21 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/06/01 00:33:21 [notice] 1#1: OS: Linux 6.5.11-linuxkit
2024/06/01 00:33:21 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/06/01 00:33:21 [notice] 1#1: start worker processes
2024/06/01 00:33:21 [notice] 1#1: start worker process 29
2024/06/01 00:33:21 [notice] 1#1: start worker process 30
```

Try visiting localhost:3005



## Starting a pod using k8s

Start a pod

```
kubectl run nginx --image=nginx --port=80
```

Check the status of the pod

```
kubectl get pods
```

Check the logs

```
kubectl logs nginx
```

```
→ week-26-prom-offline kubectl logs nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/06/01 00:32:23 [notice] 1#1: using the "epoll" event method
2024/06/01 00:32:23 [notice] 1#1: nginx/1.27.0
2024/06/01 00:32:23 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/06/01 00:32:23 [notice] 1#1: OS: Linux 6.5.11-linuxkit
2024/06/01 00:32:23 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/06/01 00:32:23 [notice] 1#1: start worker processes
2024/06/01 00:32:23 [notice] 1#1: start worker process 32
2024/06/01 00:32:23 [notice] 1#1: start worker process 33
2024/06/01 00:32:23 [notice] 1#1: start worker process 34
2024/06/01 00:32:23 [notice] 1#1: start worker process 35
2024/06/01 00:32:23 [notice] 1#1: start worker process 36
2024/06/01 00:32:23 [notice] 1#1: start worker process 37
2024/06/01 00:32:23 [notice] 1#1: start worker process 38
2024/06/01 00:32:23 [notice] 1#1: start worker process 39
```

Describe the pod to see more details

```
→ week-26-prom-offline kubectl describe pod nginx
Name:           nginx
Namespace:      default
Priority:       0
Service Account: default
Node:          local-worker/172.20.0.3 → worker 1
Start Time:    Sat, 01 Jun 2024 06:02:10 +0530
Labels:         run=nginx
Annotations:   <none>
Status:        Running
IP:            10.244.1.2
```

**What our system looks like right now**



---

Good questions to ask:-

---

How can I stop a pod?

---

How can I visit the pod? Which port is it available on?

---

How many pods can I start?

---

## Stop the pod

Stop the pod by running

```
kubectl delete pod nginx
```

Check the current state of pods

```
kubectl get pods
```

## Kubernetes manifest

---

A manifest defines the desired state for Kubernetes resources, such as Pods, Deployments, Services, etc., in a declarative manner.

---

## Original command

```
kubectl run nginx --image=nginx --port=80
```

# Manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```

## Breaking down the manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```

When was 'Pod' introduced as a k8's construct  
What we're starting  
name for the pod  
First container to start inside the pod  
Port on which the pod will listen for incoming requests

## Applying the manifest

```
kubectl apply -f manifest.yml
```

## Delete the pod

```
kubectl delete pod nginx
```

# Checkpoint

We have created

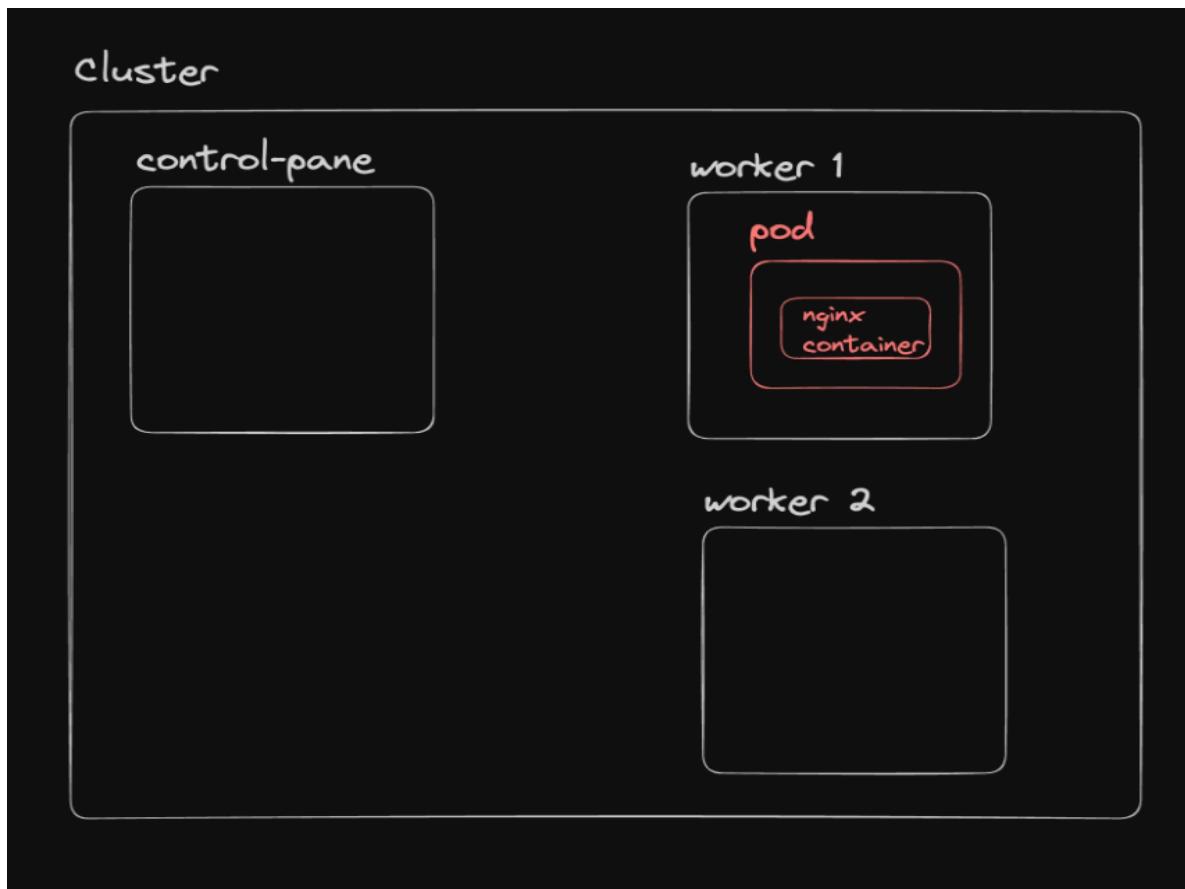
Cluster

Nodes

Images

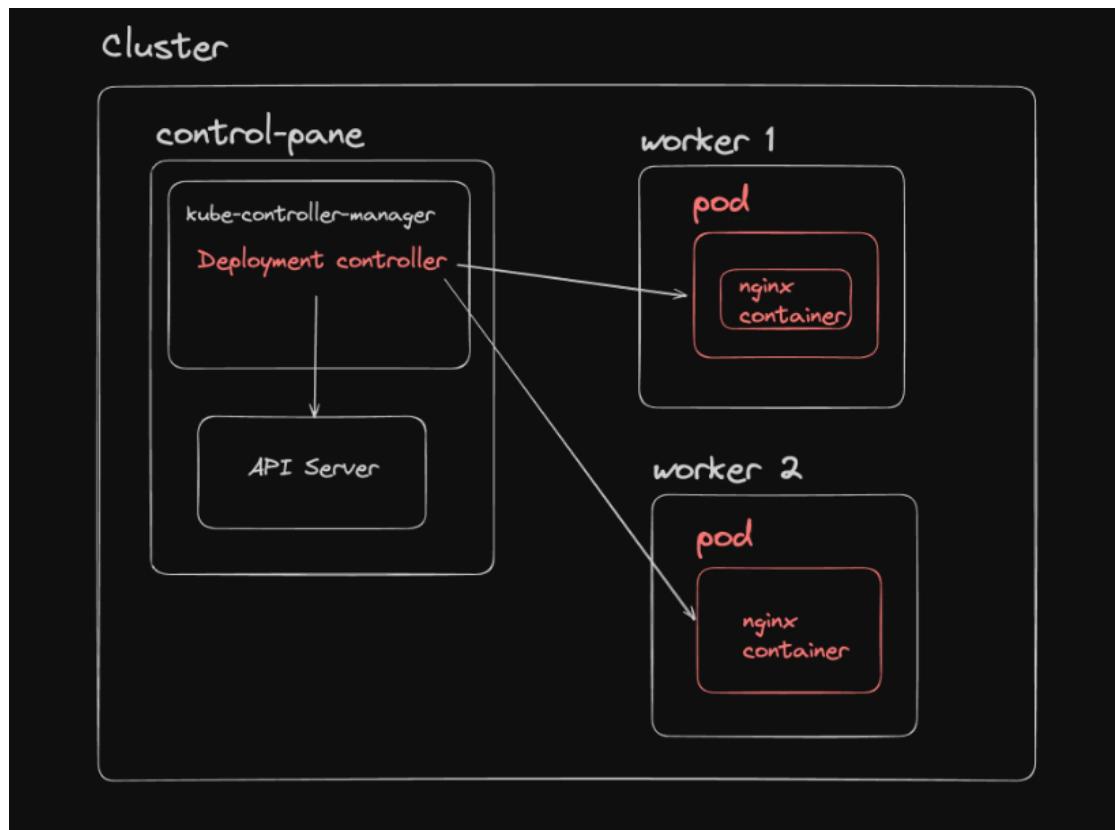
Containers

Pods



# Deployment

A Deployment in Kubernetes is a higher-level abstraction that manages a set of Pods and provides declarative updates to them. It offers features like scaling, rolling updates, and rollback capabilities, making it easier to manage the lifecycle of applications.



Pod: A Pod is the smallest and simplest Kubernetes object. It represents a single instance of a running process in your cluster, typically containing one or more containers.

Deployment: A Deployment is a higher-level controller that manages a set of identical Pods. It ensures the desired number of Pods are running and provides declarative updates to the Pods it manages.

## Key Differences Between Deployment and Pod:

### Abstraction Level:

Pod: A Pod is the smallest and simplest Kubernetes object. It represents a single instance of a running process in your cluster, typically containing one or more containers.

Deployment: A Deployment is a higher-level controller that manages a set of identical Pods. It ensures the desired number of Pods are running and provides declarative updates to the Pods it manages.

### Management:

Pod: They are ephemeral, meaning they can be created and destroyed frequently.

Deployment: Deployments manage Pods by ensuring the specified number of replicas are running at any given time. If a Pod fails, the Deployment controller replaces it automatically.

---

### Updates:

Pod: Directly updating a Pod requires manual intervention and can lead to downtime.

Deployment: Supports rolling updates, allowing you to update the Pod template (e.g., new container image) and roll out changes gradually.

If something goes wrong, you can roll back to a previous version.

---

### Scaling:

Pod: Scaling Pods manually involves creating or deleting individual Pods.

Deployment: Allows easy scaling by specifying the desired number of replicas. The Deployment controller adjusts the number of Pods automatically.

---

### Self-Healing:

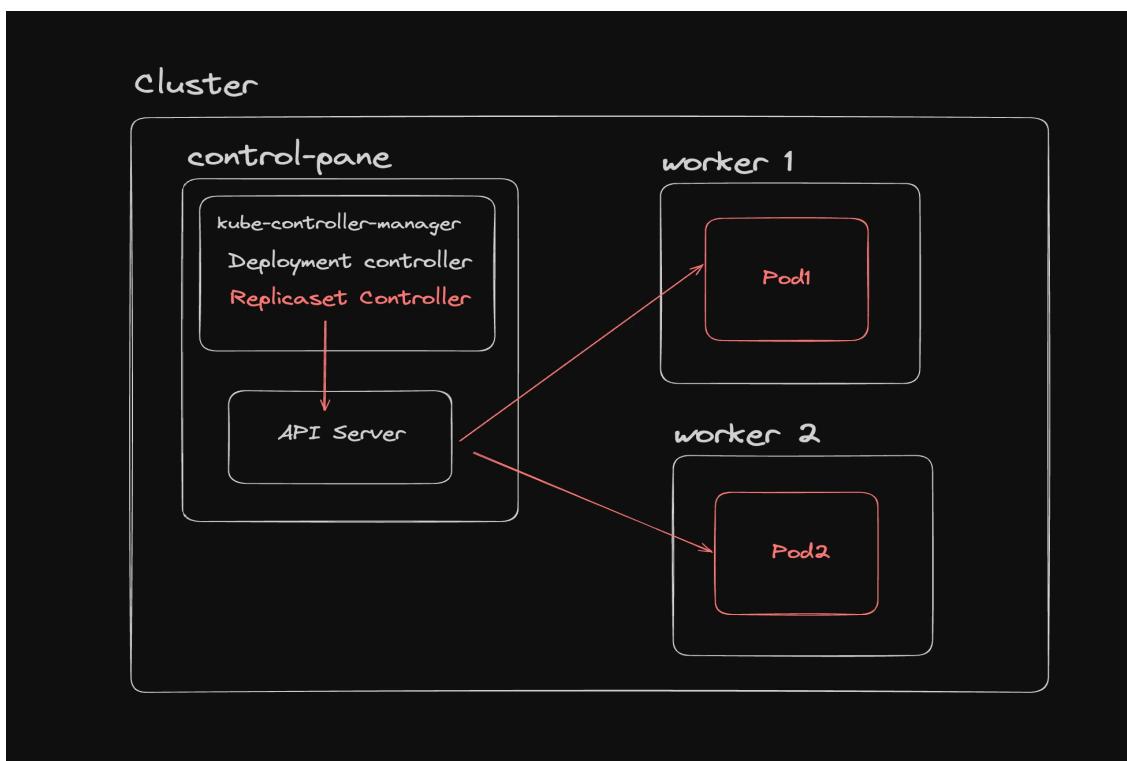
Pod: If a Pod crashes, it needs to be restarted manually unless managed by a higher-level controller like a Deployment.

Deployment: Automatically replaces failed Pods, ensuring the desired state is maintained.

# Replicaset

A ReplicaSet in Kubernetes is a controller that ensures a specified number of pod replicas are running at any given time. It is used to maintain a stable set of replica Pods running in the cluster, even if some Pods fail or are deleted.

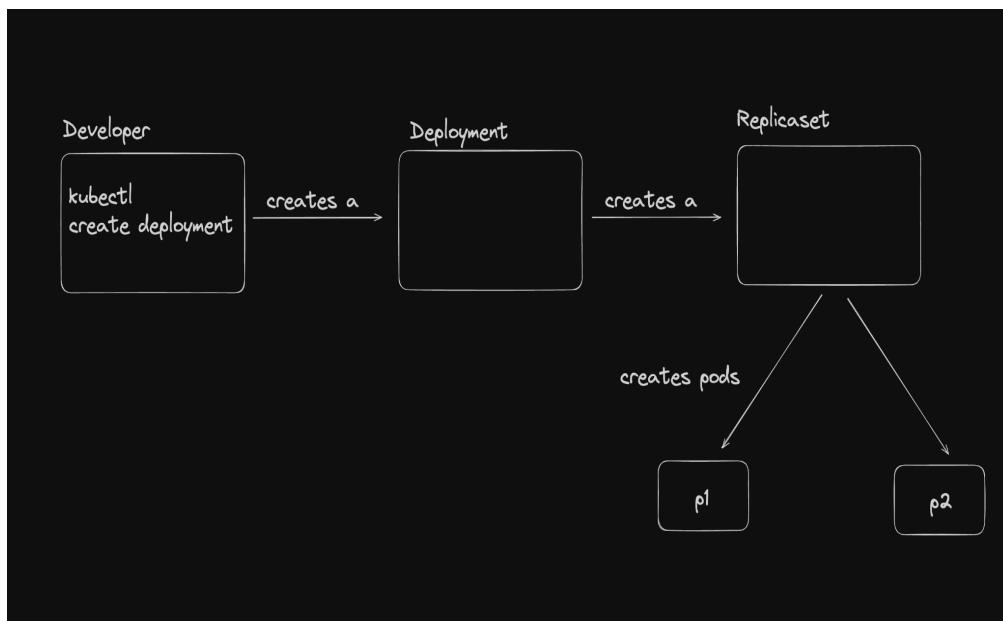
When you create a deployment, you mention the amount of replicas you want for this specific pod to run. The deployment then creates a new ReplicaSet that is responsible for creating X number of pods.



## Series of events

User creates a deployment which creates a replicaset which creates pods

If pods go down, replicaset controller ensures to bring them back up



# **Series of events**

When you run the following command, a bunch of things happen

```
kubectl create deployment \
nginx-deployment \
--image=nginx \
--port=80 \
--replicas=3
```

## Step-by-Step Breakdown:-

### 1. Command Execution:

You execute the command on a machine with kubectl installed and configured to interact with your Kubernetes cluster.

### 2. API Request:

kubectl sends a request to the Kubernetes API server to create a Deployment resource with the specified parameters.

### 3. API Server Processing:

The API server receives the request, validates it, and then processes it. If the request is valid, the API server updates the desired state of the cluster stored in etcd. The desired state now includes the new Deployment resource.

### 4. Storage in etcd:

The Deployment definition is stored in etcd, the distributed key-value store used by Kubernetes to store all its configuration data and cluster state. etcd is the source of truth for the cluster's desired state.

## 5. Deployment Controller Monitoring:

The Deployment controller, which is part of the kube-controller-manager, continuously watches the API server for changes to Deployments. It detects the new Deployment you created.

## 6. ReplicaSet Creation:

The Deployment controller creates a ReplicaSet based on the Deployment's specification. The ReplicaSet is responsible for maintaining a stable set of replica Pods running at any given time.

## 7. Pod Creation:

The ReplicaSet controller (another part of the kube-controller-manager) ensures that the desired number of Pods (in this case, 3) are created and running. It sends requests to the API server to create these Pods.

## 8. Scheduler Assignment:

The Kubernetes scheduler watches for new Pods that are in the "Pending" state. It assigns these Pods to suitable nodes in the cluster based on available resources and scheduling policies.

## 9. Node and Kubelet:

The kubelet on the selected nodes receives the Pod specifications from the API server. It then pulls the necessary container images (nginx in this case) and starts the containers

### Hierarchical Relationship in Kubernetes:-

1. Deployment (High-Level Manager): Manages application lifecycle, updates, scaling, and rollbacks by creating and managing ReplicaSets.

2. ReplicaSet (Mid-Level Manager): Ensures a specified number of Pods run, maintaining their desired state using label selectors.

3. Pod (Lowest-Level Unit): The smallest Kubernetes object, representing a running process with one or more containers.

## **Create a replicaset**

Let's not worry about deployments, lets just create a replicaset that starts 3 pods

Create rs.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:latest
      ports:
      - containerPort: 80
```

Apply the manifest

```
kubectl apply -f rs.yaml
```

Get the rs details

```
kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-replicaset	3	3	3	23s

Check the pods

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replicaset-7zp2v	1/1	Running	0	35s
nginx-replicaset-q264f	1/1	Running	0	35s
nginx-replicaset-vj42z	1/1	Running	0	35s

Try deleting a pod and check if it self heals

```
kubectl delete pod nginx-replicaset-7zp2v  
kubectl get pods
```

Try adding a pod with the `app=nginx`

```
kubectl run nginx-pod --image=nginx --labels="app=nginx"
```

Delete the replicaset

```
kubectl delete rs nginx-deployment-576c6b7b6
```

## Create a deployment

Create deployment.yml

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: nginx-deployment  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: nginx  
  template:  
    metadata:  
      labels:  
        app: nginx  
    spec:  
      containers:  
      - name: nginx  
        image: nginx:latest  
        ports:  
        - containerPort: 80
```

Apply the deployment

```
kubectl apply -f deployment.yml
```

Get the deployment

```
kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	3/3	3	3	18s

Get the rs

```
kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-deployment-576c6b7b6	3	3	3	34s

Get the pod

```
kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-576c6b7b6-b6kgk	1/1	Running	0	46s
nginx-deployment-576c6b7b6-m8ttl	1/1	Running	0	46s
nginx-deployment-576c6b7b6-n9cx4	1/1	Running	0	46s

Try deleting a pod

```
kubectl delete pod nginx-deployment-576c6b7b6-b6kgk
```

Ensure the pods are still up

```
kubectl get pods
```

Why do you need deployment?

If all that a deployment does is create a replicaset , why cant we just create rs ?

Experiment

Update the image to be nginx2 (an image that doesn't exist)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx2:latest
          ports:
            - containerPort: 80
```

Apply the new deployment

```
kubectl apply -f deployment.yml
```

Check the new **rs** now

```
kubectl get rs
NAME           DESIRED  CURRENT  READY  AGE
nginx-deployment-576c6b7b6  3        3        3     14m
nginx-deployment-5fbda799cb 1        1        0     10m
```

Check the pods

kubectl get pods					
NAME	READY	STATUS	RESTARTS	AGE	
nginx-deployment-576c6b7b6-9nlqn	1/1	Running	0	15m	
nginx-deployment-576c6b7b6-m8ttl	1/1	Running	0	16m	
nginx-deployment-576c6b7b6-n9cx4	1/1	Running	0	16m	
nginx-deployment-5fbd4799cb-fmt4f	0/1	ImagePullBackOff	0	12m	

### Role of deployment

Deployment ensures that there is a smooth deployment, and if the new image fails for some reason, the old replicaset is maintained.

Even though the rs is what does pod management, deployment is what does rs management

### **Rollbacks**

- Check the history of deployment

```
kubectl rollout history deployment/nginx-deployment
```

- Undo the last deployment

```
kubectl rollout undo deployment/nginx-deployment
```

## Create a new deployment

Replace the image to be postgres

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: postgres:latest
          ports:
            - containerPort: 80

```

- Check the new set or `rs`

```

→ kubernetes kubectl get rs
NAME           DESIRED   CURRENT   READY   AGE
nginx-deployment-576c6b7b6   3         3         3       21m
nginx-deployment-5fdbd4799cb 0         0         0       17m
nginx-deployment-7cdb767447  1         1         0       4s

→ kubernetes kubectl get pods
NAME           READY   STATUS           RESTARTS   AGE
nginx-deployment-576c6b7b6-9nlnq 1/1    Running          0          20m
nginx-deployment-576c6b7b6-m8ttl 1/1    Running          0          21m
nginx-deployment-576c6b7b6-n9cx4 1/1    Running          0          21m
nginx-deployment-7cdb767447-4d5dr 0/1    ContainerCreating 0          7s

```

- Check the pods

```

→ kubernetes kubectl get pods
NAME           READY   STATUS           RESTARTS   AGE
nginx-deployment-576c6b7b6-9nlnq 1/1    Running          0          20m
nginx-deployment-576c6b7b6-m8ttl 1/1    Running          0          21m
nginx-deployment-576c6b7b6-n9cx4 1/1    Running          0          21m
nginx-deployment-7cdb767447-4d5dr 0/1    ContainerCreating 0          7s

```

- Check pods after some time

```
→ kubernetes kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-576c6b7b6-9nlq	1/1	Running	0	23m
nginx-deployment-576c6b7b6-m8ttl	1/1	Running	0	24m
nginx-deployment-576c6b7b6-n9cx4	1/1	Running	0	24m
nginx-deployment-7cdb767447-4d5dr	0/1	CrashLoopBackOff	4 (68s ago)	3m8s

- Update the manifest to pass `POSTGRES_PASSWORD`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: postgres:latest
          ports:
            - containerPort: 80
          env:
            - name: POSTGRES_PASSWORD
              value: "yourpassword"
```

- Check pods now

```
kubectl get pods
```

```
→ kubernetes kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-576c6b7b6-9nlq	1/1	Running	0	25m
nginx-deployment-576c6b7b6-m8ttl	1/1	Running	0	26m
nginx-deployment-576c6b7b6-n9cx4	1/1	Running	0	26m
nginx-deployment-58ff5599d8-hq9r2	0/1	ContainerCreating	0	14s

- Try after some time

```
→ kubernetes kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-58ff5599d8-hq9r2	1/1	Running	0	44s
nginx-deployment-58ff5599d8-n4zxh	1/1	Running	0	25s
nginx-deployment-58ff5599d8-vbpv9	1/1	Running	0	22s

Check the rs

```
→ kubernetes kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-deployment-576c6b7b6	0	0	0	26m
nginx-deployment-58ff5599d8	3	3	3	57s
nginx-deployment-5fdbd4799cb	0	0	0	22m
nginx-deployment-7cdbb767447	0	0	0	5m38s

## How to expose the app?

Let's delete all resources and restart a `deployment` for `nginx` with 3 replicas

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

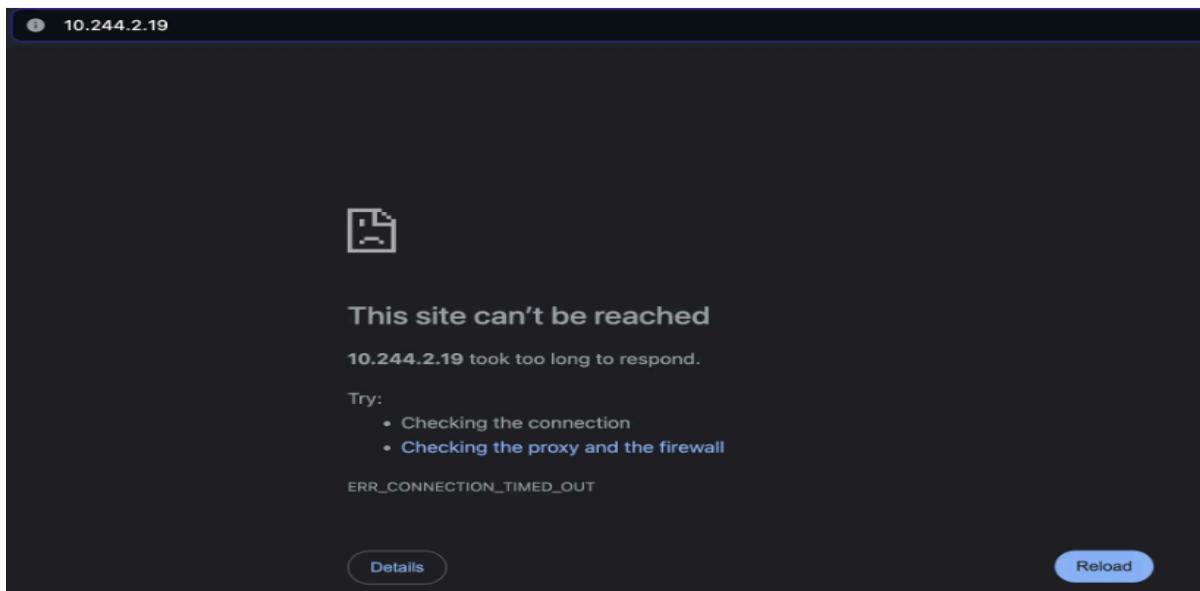
- Apply the configuration

```
kubectl apply -f deployment.yml
```

- Get all pods

```
kubectl get pods -owide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE
nginx-deployment-576c6b7b6-7jrn5  1/1    Running   0        2m19s   10.244.2.1
nginx-deployment-576c6b7b6-88fhk  1/1    Running   0        2m22s   10.244.1.
nginx-deployment-576c6b7b6-zf8ff  1/1    Running   0        2m25s   10.244.2.1
```

The IPs that you see are **private IPs**. You wont be able to access the app on it



## Services

In Kubernetes, a "Service" is an abstraction that defines a logical set of Pods and a policy by which to access them. Kubernetes Services provide a way to expose applications running on a set of Pods as network services. Here are the key points about Services in

- Create service.yml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30007 # This port can be any valid port within the NodePort range
      type: NodePort
```

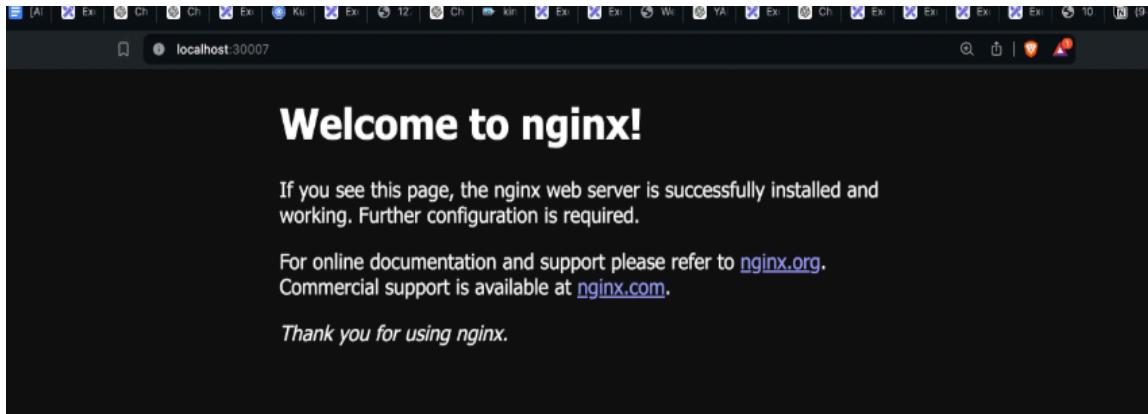
- Restart the cluster with a few extra ports exposed (create kind.yml)

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
  extraPortMappings:
    - containerPort: 30007
      hostPort: 30007
    - role: worker
    - role: worker
```

kind create cluster --config kind.yml

Re apply the deployment and the service

Visit localhost:30007



Types of services

ClusterIP

NodePort

Loadbalancer

## Loadbalancer service

In Kubernetes, a LoadBalancer service type is a way to expose a

service to external clients. When you create a Service of type

LoadBalancer, Kubernetes will automatically provision an external

load balancer from your cloud provider (e.g., AWS, Google Cloud,

Azure) to route traffic to your Kubernetes service

# Creating a kubernetes cluster in vultr

Name	Location	Created	Status	...
VKE Cluster week-27	Mumbai	2024-06-02T11:17:39+00:00	Running	***

## Create deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Apply deployment

```
kubectl apply -f deployment.yml
```

Create `service-lb.yml`

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

Apply the service

```
kubectl apply -f service-lb.yml
```

## Series of events

### Step 1 - Create your cluster

Create `kind.yml`

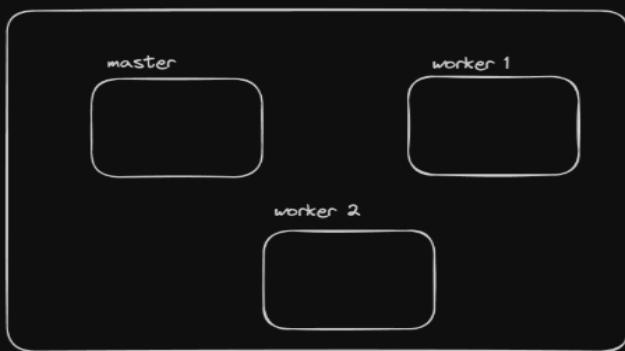
```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
  extraPortMappings:
    - containerPort: 30007
      hostPort: 30007
- role: worker
  extraPortMappings:
    - containerPort: 30007
      hostPort: 30008
- role: worker
```

- Create cluster

```
kind create cluster --config kind.yaml --name local
```

Step 1  
Creating a cluster  
Creating an image for your backend

hub.docker.com/images/nginx  
or  
hub.docker.com/images/fb\_backend



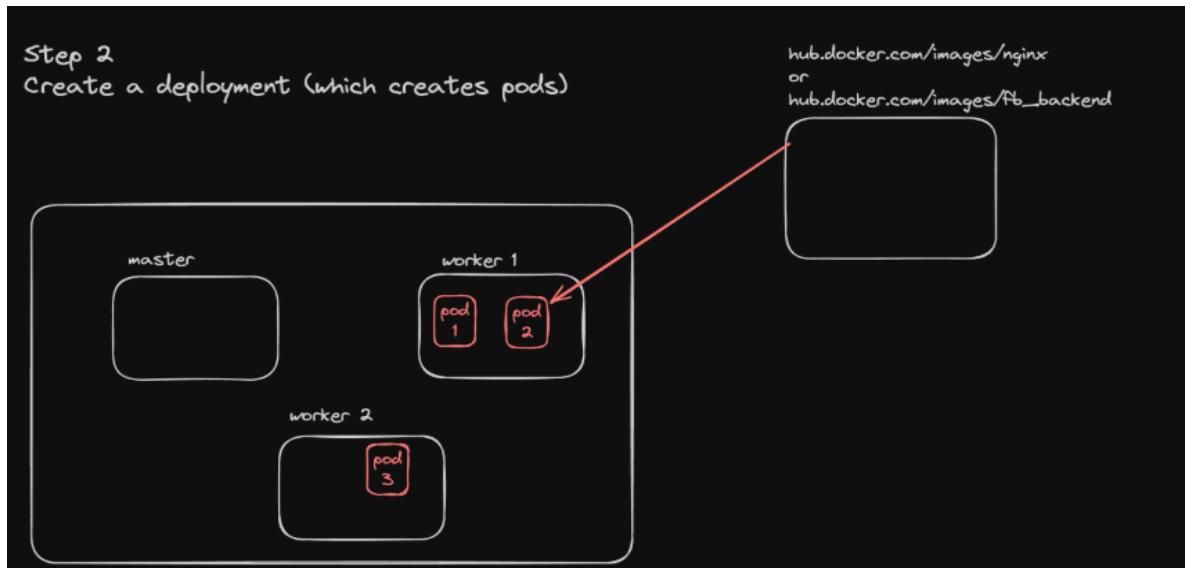
## Step 2 - Deploy your pod

Create [deployment.yaml](#)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
```

## Apply the deployment

```
kubectl apply -f deployment.yml
```



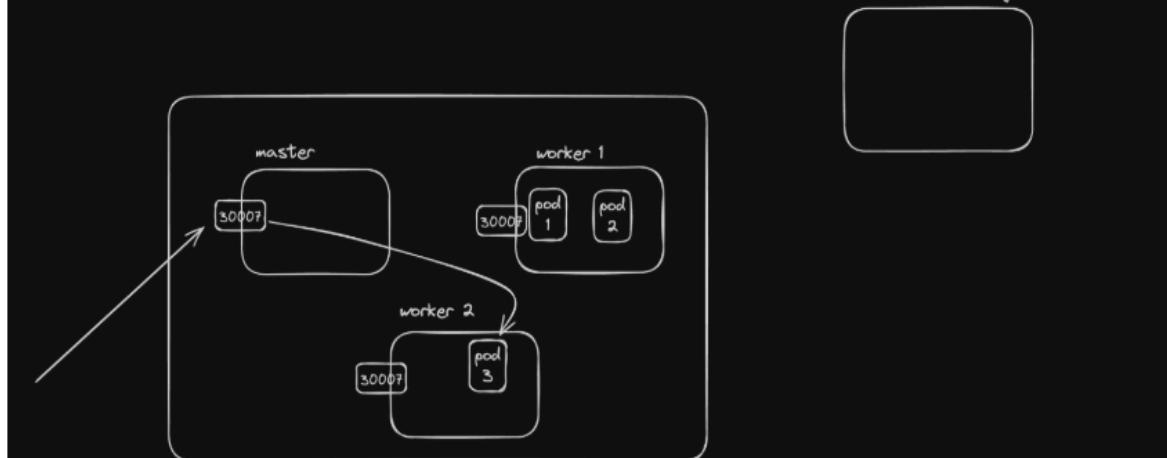
## Step 3 - Expose your app over a NodePort

Create service.yml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30007 # This port can be any valid port within the NodePort range
  type: NodePort
```

**Step 3**  
Expose a service over NodePort

hub.docker.com/images/nginx  
or  
hub.docker.com/images/fb\_backend



## Step 4 - Expose it over a LoadBalancer

Create a load balancer service (service-lb.yml)

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

```
kubectl apply service-lb.yml
```



## Check the cloud dashboard

The screenshot shows a cloud dashboard interface. On the left, there is a sidebar with a tree view of services: Products (Compute, Cloud Storage, Kubernetes, Serverless, Container Registry, Databases), Load Balancers (selected), CDN, Network, and Orchestration. A red arrow points from the "Load Balancers" section in the sidebar to the list of load balancers on the right. The main area displays a table titled "Load Balancers" with the following data:

Name	Location	Charges	Status
a1a84c241140c4db7a04d5e9f03015ea 5cb65774-53a1-4ed9-9624-35d178403117	Mumbai	\$0.02	Running

Below the table, there is a "Frequently Asked Questions" section with a blue plus sign icon.